

Declaring Fortran Device Data

- Variables / arrays with device attribute are allocated in device memory

real, device, allocatable :: a(:)

real, allocatable :: a(:)

attributes(device) :: a

- In a host subroutine or function
 - device allocatables and automatics may be declared
 - device variables and arrays may be passed to other host subroutines or functions (explicit interface)
 - device variables and arrays may be passed to kernel subroutines

Declaring Fortran Module Data

- Variables / arrays with device attribute are allocated in device memory

```
module mm
  real, device, allocatable :: a(:)
  real, device :: x, y(10)
  real, constant :: c1, c2(10)
  integer, device :: n
contains
  attributes(global) subroutine s( b )
  ...
```

- Module data must be fixed size, or allocatable

Allocating Data

- Fortran allocate / deallocate statement

```
real, device, allocatable :: a(:, :), b
allocate( a(1:n,1:m), b )
....
deallocate( a, b )
```

- arrays or variables with device attribute are allocated in device memory

- Allocate is done by the host subprogram
- Memory is not virtual, you can run out
- Device memory is shared among users / processes, you can have deadlock
- STAT=i var clause to catch and test for errors

Copying Data to / from Device

- Assignment statements

```
real, device, allocatable :: a(:, :), b
allocate( a(1:n,1:m), b )
a(1:n,1:m) = x(1:n,1:m) ! copies to device
b = 99.0
....
x(1:n,1:m) = a(1:n,1:m) ! copies from device
y = b
deallocate( a, b )
```

- Data copy may be noncontiguous, but will then be slower (multiple DMAs)
- Data copy to / from host pinned memory will be faster
- Asynchronous copies currently require API interface

Launching Kernels

- Subroutine call with chevron syntax for launch configuration

```
call vaddkernel <<<(N+31)/32,32 >>> (A,B,C,N)
```

```
type(dim3) :: g, b
```

```
g = dim3((N+31)/32, 1, 1)
```

```
b = dim3( 32, 1, 1 )
```

```
call vaddkernel <<< g, b >>> ( A, B, C, N )
```

- Interface must be explicit
 - In the same module as the host subprogram
 - In a module that the host subprogram uses
 - Declared in an interface block
- The launch is asynchronous
 - host program continues, may issue other launches

CUDA Errors

- Out of memory
- Launch failure (array out of bounds, ...)
- No device found
- Invalid device code (compute capability mismatch)

Test for error:

```
ir = cudaGetLastError()  
if( ir ) print *, cudaGetErrorString( ir )
```

```
ir = cudaGetLastError();  
if( ir ) printf( "%s\n", cudaGetErrorString(ir) );
```


Writing a CUDA Kernel (1)

- C: global attribute on the function header, must be void type
 - `__global__ void kernel (...){...}`
- F: global attribute on the subroutine statement
 - `attributes(global) subroutine kernel (A, B, C, N)`
- May declare scalars, fixed size arrays in local memory
- May declare shared memory arrays
 - C: `__shared__ float sm(16,16);`
 - F: `real, shared :: sm(16,16)`
 - Limited amount of shared memory available (16KB, 48KB)
 - shared among all threads in the same thread block
- Data types allowed
 - int (long,short,char), float, double, struct, union, ...
 - integer(1,2,4,8), logical(1,2,4,8), real(4,8), complex(4,8), derivedtype

Writing a CUDA Kernel (2)

- Predefined variables
 - `blockIdx`, `threadIdx`, `gridDim`, `blockDim`, `warpSize`
- Executable statements in a kernel
 - assignment
 - for, do, while, if, goto, switch
 - function call to device function
 - intrinsic function call
 - most intrinsics implemented in header files

Writing a CUDA Kernel (3)

- Fortran disallowed statements include
 - read, write, print, open, close, inquire, format, other IO except now some limited support for list-directed (print *)
 - allocate, deallocate, until PGI 13.0
 - Fortran pointer assignment
 - recursive procedure calls, direct or indirect
 - ENTRY statement, optional arguments, alternate return
 - data initialization, SAVED data
 - assigned goto, ASSIGN statement
 - stop, pause

Exercise Solution

- `multidim_solution.cuf`
- In the book on page 12

Kernel Loop Directives (CUF Kernels)

- Automatic kernel generation and invocation of host code region (arrays used in loops must reside on GPU)

```
program incTest
  use cudafor
  implicit none
  integer, parameter :: n = 256
  integer :: a(n), b
  integer, device :: a_d(n)

  a = 1; b = 3; a_d = a

  !$cuf kernel do <<<*,*>>>
  do i = 1, n
    a_d(i) = a_d(i)+b
  enddo

  a = a_d
  if (all(a == 4)) write(*,*) 'Test Passed'
end program incTest
```

Kernel Loop Directives (CUF Kernels)

- Multidimensional arrays

```
!$cuf kernel do(2) <<< *,* >>>  
do j = 1, ny  
  do i = 1, nx  
    a_d(i,j) = b_d(i,j) + c_d(i,j)  
  enddo  
enddo
```

- Can specify parts of execution parameter

```
!$cuf kernel do(2) <<<(*,*), (32,4)>>>
```


Kernel Loop Directives (CUF Kernels)

Syntax:

The general form of the kernel directive is:

```
!$cuf kernel do[(n)] <<< grid, block [ optional stream ] >>>
```

The compiler maps the launch configuration specified by the grid and block values onto the outermost n loops, starting at loop n and working out. The grid and block values can be an integer scalar or a parenthesized list. Alternatively, using asterisks tells the compiler to choose a thread block shape and/or compute the grid shape from the thread block shape and the loop limits. Loops which are not mapped onto the grid and block values are run sequentially on each thread.

Kernel Loop Directive Syntax

```
!$cuf kernel do(7) <<< *,(32,8,1,1,1,1,1) >>>  
  do j7 = 1, is27  
    do j6 = 1, is26  
      do j5 = 1, is25  
        do j4 = 1, is24  
          do j3 = 1, is23  
            do j2 = 1, is22  
              do j1 = 1, is21  
                . . .
```


Reduction using CUF Kernels

- Compiler recognizes use of scalar reduction and generates one result

```
rsum = 0.0
!$cuf kernel do <<<*,*>>>
do i = 1, nx
    rsum = rsum + a_d(i)
enddo
```

!\$CUF kernel directives

```
module madd_device_module
  use cudafor
  contains
    subroutine madd_dev(a,b,c,sum,n1,n2)
      real,dimension(:, :, ),device :: a,b,c
      real :: sum
      integer :: n1,n2
      type(dim3) :: grid, block
```

!\$cuf kernel do (2) <<<(*,*) , (32,4)>>>

```
      do j = 1,n2
        do i = 1,n1
          a(i,j) = b(i,j) + c(i,j)
          sum = sum + a(i,j)
        enddo
      enddo
    end subroutine
end module
```

Equivalent
hand-written
CUDA kernels

```
module madd_device_module
  use cudafor
  implicit none
  contains
    attributes(global) subroutine madd_kernel(a,b,c,blocksum,n1,n2)
      real, dimension(:, :), device :: a,b,c
      real, dimension(:) :: blocksum
      integer, value :: n1,n2
      integer :: i,j,tindex,tneighbor,bindex
      real :: mysum
      real, shared :: bsum(256)
      ! Do this thread's work
      mysum = 0.0
      do j = threadidx%y + (blockidx%y-1)*blockdim%y, n2, blockdim%y*griddim%y
        do i = threadidx%x + (blockidx%x-1)*blockdim%x, n1, blockdim%x*griddim%x
          a(i,j) = b(i,j) + c(i,j)
          mysum = mysum + a(i,j) ! accumulates partial sum per thread
        enddo
      enddo
      ! Now add up all partial sums for the whole thread block
      ! Compute this thread's linear index in the thread block
      ! We assume 256 threads in the thread block
      tindex = threadidx%x + (threadidx%y-1)*blockdim%x
      ! Store this thread's partial sum in the shared memory block
      bsum(tindex) = mysum
      call syncthreads()
      ! Accumulate all the partial sums for this thread block to a single value
      tneighbor = 128
      do while( tneighbor >= 1 )
        if( tindex <= tneighbor ) &
          bsum(tindex) = bsum(tindex) + bsum(tindex+tnneighbor)
        tneighbor = tneighbor / 2
        call syncthreads()
      enddo
      ! Store the partial sum for the thread block
      bindex = blockidx%x + (blockidx%y-1)*griddim%x
      if( tindex == 1 ) blocksum(bindex) = bsum(1)
    end subroutine

    ! Add up partial sums for all thread blocks to a single cumulative sum
    attributes(global) subroutine madd_sum_kernel(blocksum,dsum,nb)
      real, dimension(:) :: blocksum
      real :: dsum
      integer, value :: nb
      real, shared :: bsum(256)
      integer :: tindex,tnneighbor,i
      ! Again, we assume 256 threads in the thread block
      ! accumulate a partial sum for each thread
      tindex = threadidx%x
      bsum(tindex) = 0.0
      do i = tindex, nb, blockdim%x
        bsum(tindex) = bsum(tindex) + blocksum(i)
      enddo
      call syncthreads()
      ! This code is copied from the previous kernel
      ! Accumulate all the partial sums for this thread block to a single value
      ! Since there is only one thread block, this single value is the final result
      tneighbor = 128
      do while( tneighbor >= 1 )
        if( tindex <= tneighbor ) &
          bsum(tindex) = bsum(tindex) + bsum(tindex+tnneighbor)
        tneighbor = tneighbor / 2
        call syncthreads()
      enddo
      if( tindex == 1 ) dsum = bsum(1)
    end subroutine

    subroutine madd_dev(a,b,c,dsum,n1,n2)
      real, dimension(:, :), device :: a,b,c
      real, device :: dsum
      real, dimension(:), allocatable, device :: blocksum
      integer :: n1,n2,nb
      type(dim3) :: grid, block
      integer :: i
      ! Compute grid/block size; block size must be 256 threads
      grid = dim3((n1+31)/32, (n2+7)/8, 1)
      block = dim3(32,8,1)
      nb = grid%x * grid%y
      allocate(blocksum(1:nb))
      call madd_kernel<<< grid, block >>>>(a,b,c,blocksum,n1,n2)
      call madd_sum_kernel<<< 1, 256 >>>>(blocksum,dsum,nb)
      ! = cudaThreadSynchronize() ! don't deallocate too early
      deallocate(blocksum)
    end subroutine
  end module
```


Exercise - CUF Kernels

- Modify multidim.cuf to use a CUF kernel
- Solution in multidim_CUF_solution.cuf
- Extra credit: calculate the sum of all elements in the array
- Consult section 3.7 *Kernel Loop Directives* in book if needed

Compute Capabilities

Architecture	Tesla				Fermi		Kepler			Maxwell		Pascal		
Compute capabilities	1	1.1	1.2	1.3	2	2.1	3	3.5	3.7	5.0	5.x	6.0	6.x	
Double precision					Full		Full Full					Full		
3D grids														
Max # threads per block	512				1024									
Shared memory per MP	16Kb						48Kb (16/48,48/16)			48Kb +(32/32)				
.....							

- All these values are returned by `cudaGetDeviceProperties` or `pgaccelinfo`
- Target GPU can be specified with `-Mcuda=ccxy`

Device Query Code

```
type (cudaDeviceProp) :: prop
integer :: nDevices=0, i, ierr

ierr = cudaGetDeviceCount(nDevices)
...
do i = 0, nDevices-1

    write(*, "('Device Number: ',i0)") i
    ierr = cudaGetDeviceProperties(prop, i)

    write(*, "('  Device Name: ',a)") trim(prop%name)
    write(*, "('  Compute Capability: ',i0,'.',i0)") &
        prop%major, prop%minor
    write(*, "('  Number of Multiprocessors: ',i0)") &
        prop%multiProcessorCount
    write(*, "('  Max Threads per Multiprocessor: ',i0)") &
        prop%maxThreadsPerMultiprocessor
    write(*, "('  Global Memory (GB): ',f9.3,/)") &
        prop%totalGlobalMem/1024.0**3
    ...
```

Compile and Run deviceQuery.cuf

- Note: devices are enumerated from 0
- Full blown deviceQuery? `pgaccelinfo` utility

Compilation

- PGI's Fortran compiler

- All source code with **.cuf** or **.CUF** is compiled as CUDA Fortran enabled automatically
- Flag to target architecture (eg. **-Mcuda=cc35**)
 - **-Mcuda=emu** specifies emulation mode
- Flag to target CUDA Toolkit version (eg. **-Mcuda=cuda5.5**)
 - **-Mcuda=fastmath** enables faster intrinsics (**__sinf()**)
 - **-Mcuda=nofma** turns off fused multiply-add
 - **-Mcuda=maxregcount:<n>** limits register usage per thread
 - **-Mcuda=ptxinfo** prints memory usage per kernel

Use **pgf90 -Mcuda -help** for a full list